

TheTowerDNN and GraphRanker were both trained using contrastive learning and HPO objective metrics that were partitioned into head, torso and tail movies (tier=0,1,2, resp) as defined by the training positive dataset movie ratings frequency.

- The movie embeddings generated from the TwoTowerDNN Candidate model show decent separation of clusters in the t-SNE and UMAP plots with spatially overlapping tier=0,1,2. They look healthy
- The train, val learning curves for `ndcg_<tier>` in the TwoTowerDNN and GraphRanker training all show healthy learning above random chance metrics.
- The batch sizes have been adjusted to `>= 1024` in order to be able to learn the tail distribution. Binomial distribution and Poisson distribution were to calculate the batch sizes. See `recommender_systems/docs/mlops/movie_tier_counts.txt` (add github ref).
- End-to-end inference metrics using the trained models and the test dataset:
 - 48_285 test dataset ratings
 - 3_883 num_catalog_movies
 - "ndcg_head_20": 0.044804782414759374,
 - "ndcg_torso_20": 0.02441173223123198,
 - "ndcg_tail_20": 0.0060162138110536065,
 - "random_ndcg_head_20": 0.003064348921321239,
 - "random_ndcg_torso_20": 0.0028287642332191837,
 - "random_ndcg_tail_20": 0.0020430092333822906,
 - "recall_head_20": 0.0566528978866582,
 - "recall_torso_20": 0.052628152590776876
 - "recall_tail_20": 0.018960269097068005,
 - "random_recall_head_20": 0.0051841404718014975,
 - "random_recall_torso_20": 0.005218351778443615,
 - "random_recall_tail_20": 0.005150656708730404,
 - "catalog coverage": 0.7334535121917725,
 - "count_users_head_20": 4584,
 - "count_users_torso_20": 4063,
 - "count_users_tail_20": 731,
 - "count_movie_catalog_head": 565,
 - "count_movie_catalog_torso": 1709,
 - "count_movie_catalog_tail": 1609,
 - "count_test_head": 565,
 - "count_test_torso": 1692,
 - "count_test_tail": 470,
 - "frac_pred_tiers_head_20": 0.37626527050610753,
 - "frac_pred_tiers_torso_20": 0.5706128476495198,
 - "frac_pred_tiers_tail_20": 0.1140902872777015
- =>
count_movie=np.array([565, 1709, 1609])
count_test=np.array([565, 1692, 470])
count_users=np.array([4584, 4063, 731])
count_movie/count_movie.sum() = array([0.1455, 0.4401, 0.4144])
count_users/count_users.sum() =array([0.4888, 0.4332, 0.0780])
count_test/count_test.sum() = array([0.2072, 0.6205, 0.1724])
- =>
 - "frac_pred_tiers_tail_20": 0.11409

- $0.114 * k=20$ yields 2.28 niche (tail) movies" for users who rated tier=2 movies in test ("niche users")
- This is above 0 so is predicting niche items to niche users, but is it different than random selection? Yes, uniform random selection is $\text{count_movie}/\text{count_movie.sum()} = \text{array}([0.1455, 0.4401, 0.4144])$. So 20 uniform random recommendations * $[0.1455, 0.4401, 0.4144] = [2.9, 8.8, 8.3]$.
The $\text{frac_pred_tiers} = 20 * [0.376, 0.571, 0.114] = [7.5, 11.4, 2.3]$ which is clearly not the uniform random distribution. So this is more evidence that the end-to-end inference is able to make niche recommendations.
- =>
 - The $\text{ndcg_}<\text{tier}>$ and $\text{recal_}<\text{tier}>$ metrics are all 3 to 10 times the metric expected from random selection, which shows that the models successfully predict niche movies for niche users while still predicting popular movies to popular users.
- =>
 - "catalog coverage": 0.73
The recommendation pipeline is recommending to users, far more than the small number of very popular movies, hence is overcoming the popular movie bias.

Could the niche user/movie recommendations be improved by increasing num_candidates to be 10x to 100x top_k?

- I took the intersection of users who rated tier=2 movies in the train, val, and positive datasets, which is 222 users. For those users, ANN searches w.r.t. each dataset (using their first timestamp in the dataset) were made to get the top k number of recommendations. The average number of movie tiers in each set of top k recommendations was determined. Here are those averages:
dataset_k-----avg counts in the k----- avg counts / k
- | | | |
|-------------|---------|---------|
| test_00070: | 21.432 | : 0.306 |
| test_00100: | 31.914 | : 0.319 |
| test_00200: | 69.311 | : 0.347 |
| test_00300: | 108.378 | : 0.361 |
| test_00400: | 149.604 | : 0.374 |
| test_00500: | 190.964 | : 0.382 |
| test_00600: | 232.829 | : 0.388 |
| test_00700: | 275.928 | : 0.394 |
| test_00800: | 318.568 | : 0.398 |
| test_00900: | 362.009 | : 0.402 |
| test_01000: | 405.973 | : 0.406 |
| test_01100: | 449.500 | : 0.409 |
| test_01200: | 492.847 | : 0.411 |
| test_01300: | 535.955 | : 0.412 |
| test_01400: | 579.090 | : 0.414 |
| test_01500: | 622.946 | : 0.415 |
| test_01600: | 666.887 | : 0.417 |
| test_01700: | 711.126 | : 0.418 |
| test_01800: | 755.176 | : 0.420 |
| test_01900: | 800.099 | : 0.421 |
| test_02000: | 844.378 | : 0.422 |

train_00070: 20.802 : 0.297
 train_00100: 31.104 : 0.311
 train_00200: 67.811 : 0.339
 train_00300: 105.662 : 0.352
 train_00400: 145.113 : 0.363
 train_00500: 185.563 : 0.371
 train_00600: 226.644 : 0.378
 train_00700: 268.437 : 0.383
 train_00800: 310.568 : 0.388
 train_00900: 353.482 : 0.393
 train_01000: 396.707 : 0.397
 train_01100: 439.568 : 0.400
 train_01200: 483.378 : 0.403
 train_01300: 526.806 : 0.405
 train_01400: 570.216 : 0.407
 train_01500: 614.586 : 0.410
 train_01600: 658.216 : 0.411
 train_01700: 702.824 : 0.413
 train_01800: 747.311 : 0.415
 train_01900: 792.032 : 0.417
 train_02000: 836.432 : 0.418
 val_00070: 21.095 : 0.301
 val_00100: 31.500 : 0.315
 val_00200: 68.374 : 0.342
 val_00300: 106.865 : 0.356
 val_00400: 147.279 : 0.368
 val_00500: 188.306 : 0.377
 val_00600: 229.671 : 0.383
 val_00700: 272.131 : 0.389
 val_00800: 314.671 : 0.393
 val_00900: 357.806 : 0.398
 val_01000: 401.464 : 0.401
 val_01100: 445.036 : 0.405
 val_01200: 488.315 : 0.407
 val_01300: 531.275 : 0.409
 val_01400: 574.622 : 0.410
 val_01500: 618.757 : 0.413
 val_01600: 662.662 : 0.414
 val_01700: 706.874 : 0.416
 val_01800: 751.513 : 0.418
 val_01900: 796.013 : 0.419
 val_02000: 840.189 : 0.420

- ⇒ this shows that increasing the number of candidates would increase the fraction of tier=2 movies in the recommendations.
 - Increasing num_candidates increases the graph size by by factor if candidate increase. E.g. the metrics at top of file are from num_candidates=70 model. If we increase the number of candidates to 2000 the factor of increase is nearly 30.

- In the GraphRanker, the graph of senders and receivers given to the GATV2 uses a star bipartite morphology centered around the user node and so the memory scaling is a linear $O(c)$ so the VRAM needed will scale linearly too.
- The VRAM consumed for the num_candidates=70 model was roughly 900 MiB. so scaling up to 200 candidates would use about 2.6GiB of the available 30 GB RAM.
- The runtime complexity is $O(E+V)$ and over a batch size of 1024 a model training using num_candidates=200 would be roughly 2-3 times longer than a model trained with num_candidates=70.
- num_candidates=70 was enough to implement niche user recommendations.
 - Increasing num_candidates to 2000 would give a moderate increase in the fraction of tier=2 recommendations returned to niche users, but at the cost of nearly 20 times the elapsed runtime.

$$(1024 * (50 + 2000)) / (1024 * (50 + 70))$$